



OLTP workload mapping for database tuning

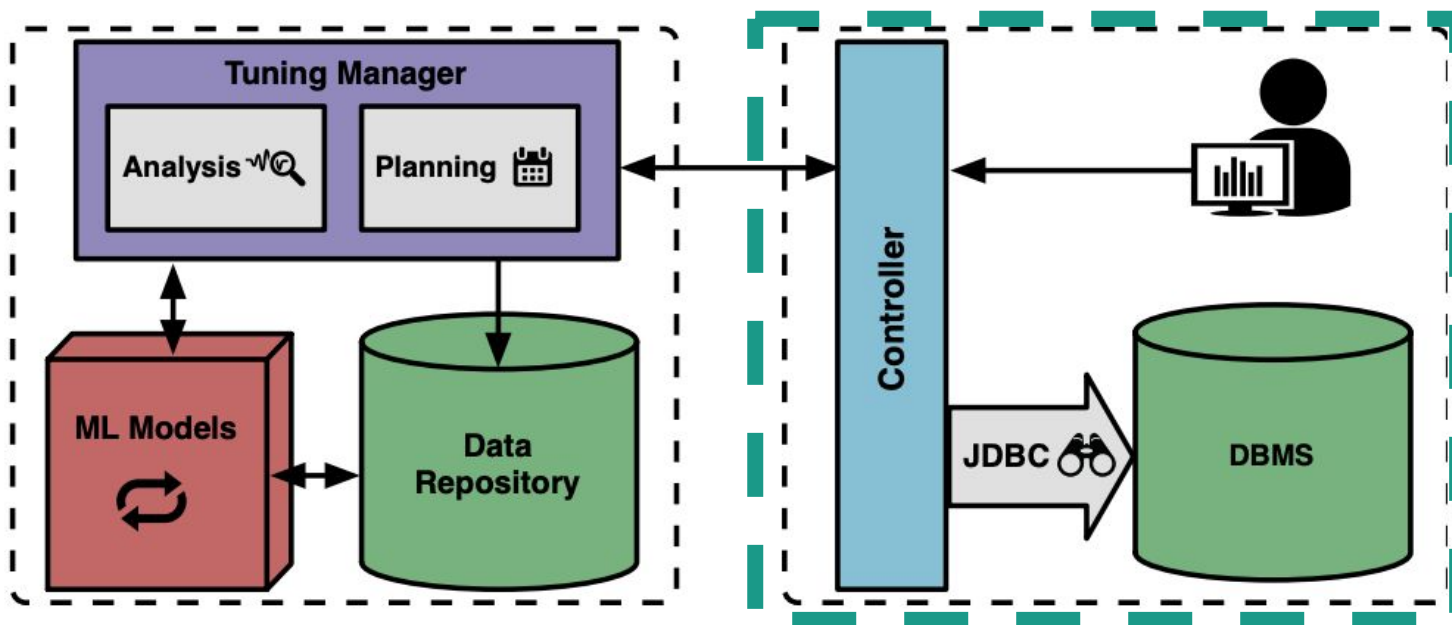
CS310 - Valentin Kodderitzsch




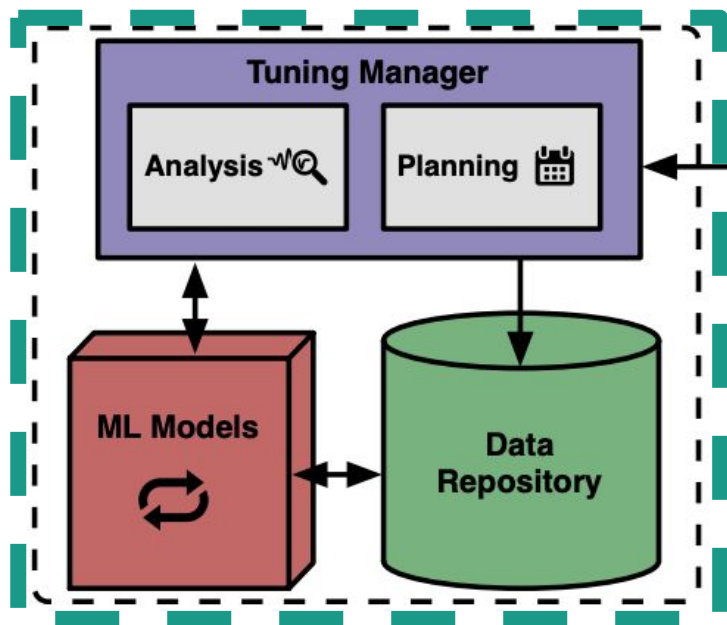
Big picture

1. Improve database performance
2. Changing database settings
3. Reduce tuning time

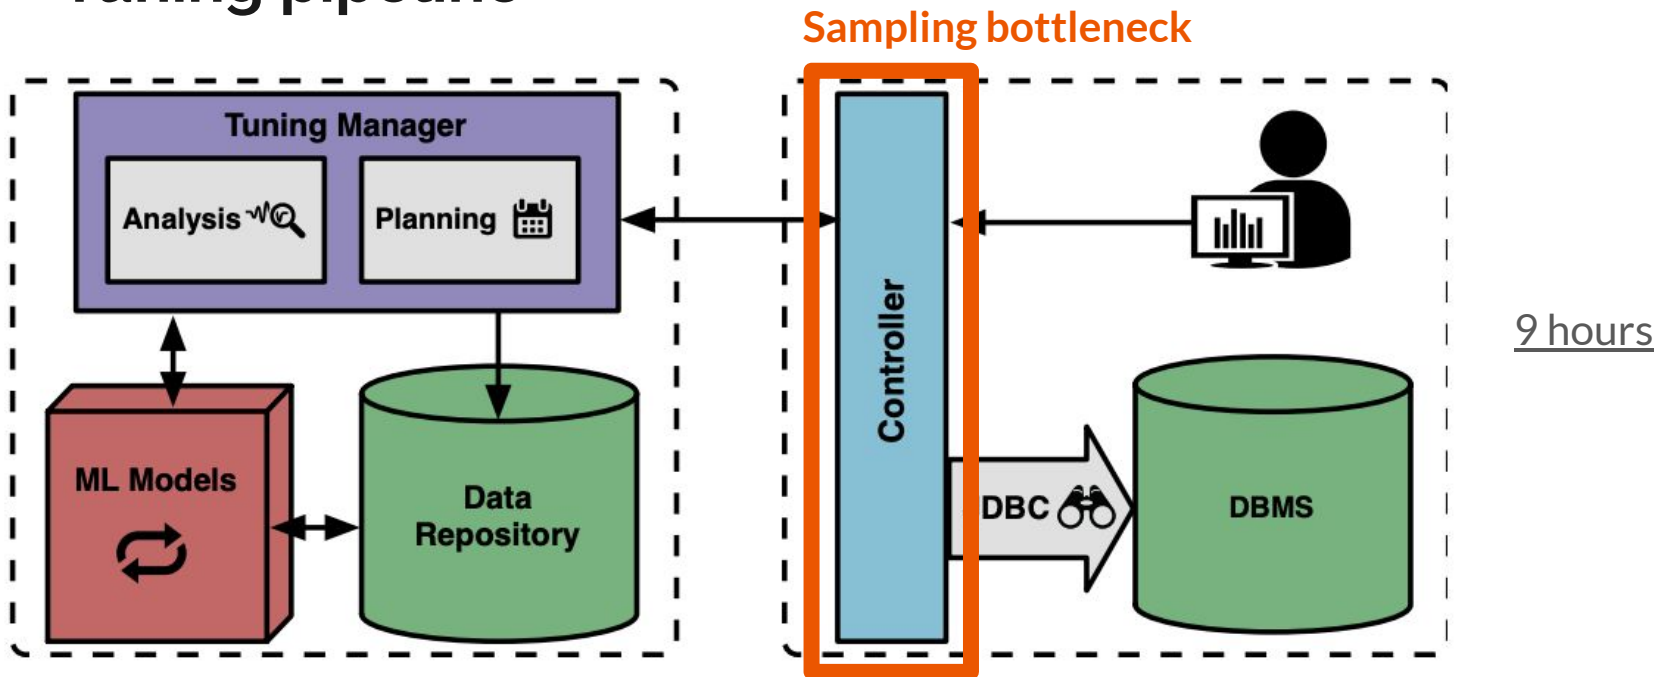
Tuning pipeline




Tuning pipeline




Tuning pipeline



What if the workload changes?



9 hours




9 hours again??

Can we reduce the sampling bottleneck by mapping between workloads?


25% faster

In only $\frac{1}{3}$ of the tuning time

Key intuition




Old




New

Reuse samples

Key intuition



Old



New

Forget old samples



New sampling function

$$n_{new} = f(n_{old}, \theta) = \theta * n_{old}$$

$$0 \leq \theta \leq 1$$

$$n_{old} = 100$$

Theta before sampling??



Jensen-Shannon divergence

Relative entropy between two probability distributions

Symmetric version of the Kullback–Leibler divergence (KLD or DKL)

Base 2 gives a score between 0 and 1

0 → Identical

1 → Completely different



Jensen-Shannon divergence

$$\text{JSD}(P \parallel Q) = \frac{1}{2}D(P \parallel M) + \frac{1}{2}D(Q \parallel M),$$

Score between 0 and 1

→ Sampling function

→ Ensemble method

↑ ↑
Prob. distributions P and Q

P and Q ??

$$M = \frac{1}{2}(P + Q)$$

$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left(\frac{P(x)}{Q(x)} \right).$$

Key assumption


$$JSD(A_{default} || B_{default}) \approx JSD(A_{all_samples} || B_{all_samples})$$

A: PDF of Default logs


1	Transaction_Type	Index	Transaction_Nom	Star
2	4	Delivery	1678361810-799424-00000-0-0	
3	1	Neuroder	1678361810-799504-7234-0-0	
4	2	Payment	1678361810-799507-54640-0-0	
5	1	Neuroder	1678361810-883172-00000-0-0	
6	1	Neuroder	1678361810-883192-57378-7-0	
7	2	Payment	1678361810-883195-5876-4-0	
8	4	Delivery	1678361810-882710-80224-0-0	
2033101	2	Payment	1678362105-798426-00000-0-0	
2033102	2	Payment	1678362105-798475-5882500-0-0	
2033103	2	Payment	1678362105-798137-2088161-0-0	
2033201	3	OrderStatus	1678362105-798137-4000001-2-0	

B: PDF of Default logs

1	Transaction_Type	Index	Transaction_Nom	Star
2	4	Delivery	1678361810-799424-00000-0-0	
3	1	Neuroder	1678361810-799504-7234-0-0	
4	2	Payment	1678361810-799507-54640-0-0	
5	1	Neuroder	1678361810-883172-00000-0-0	
6	1	Neuroder	1678361810-883192-57378-7-0	
7	2	Payment	1678361810-883195-5876-4-0	
8	4	Delivery	1678361810-882710-80224-0-0	
2033101	2	Payment	1678362105-798426-00000-0-0	
2033102	2	Payment	1678362105-798475-5882500-0-0	
2033103	2	Payment	1678362105-798137-2088161-0-0	
2033201	3	OrderStatus	1678362105-798137-4000001-2-0	



A: PDF all samples



B: PDF all samples

Additional statistics



Workload

Synthetic

ID = (45, 45, 3, 3, 4) or (20, 20, 20, 39, 1), etc.

⇒ 4.6 million

(Difallah, Pavlo, Curino; 2013)

```
26      <time>300</time>
27      <weights>30,45,5,10,10</weights>  <!-- E -->
28      </work>
29  </works>
30
31  <!-- TPCC specific -->
32  <transactiontypes>
33      <transactiontype>
34          <name>NewOrder</name>
35      </transactiontype>
36      <transactiontype>
37          <name>Payment</name>
38      </transactiontype>
39      <transactiontype>
40          <name>OrderStatus</name>
41      </transactiontype>
42      <transactiontype>
43          <name>Delivery</name>
44      </transactiontype>
45      <transactiontype>
46          <name>StockLevel</name>
47      </transactiontype>
48  </transactiontypes>
49 </parameters>
```

DB settings

Index	Name	What	Min	Default	Max (hardware dependent)
1	effective_cache_size	Memory (RAM)	8 KB	4 GB	80% of 8GB
2	maintenance_work_mem	Memory	1 MB	64 MB	80% of 8GB
3	max_wal_size	Storage (log)	32 MB	1 GB	80% of 5GB
4	max_worker_processes	CPU	0	4	8
5	shared_buffers	Memory	128 KB	8 MB	80% of 8GB
6	temp_buffers	Memory	128 KB	8 MB	80% of 8GB
7	wal_buffers	Memory	64 KB	64 KB (same as min)	80% of 8GB
8	work_mem	Memory	64 KB	4 MB	80% of 8GB

Domain for sampling algorithm

Algorithm returns n features

```
1 {
2   "effective_cache_size": 1426659,
3   "maintenance_work_mem": 4509998,
4   "max_wal_size": 1878597,
5   "max_worker_processes": 5,
6   "shared_buffers": 5149266,
7   "temp_buffers": 42646,
8   "wal_buffers": 126151,
9   "work_mem": 7625
10 }
```

Feature 1

```
1 {
2   "effective_cache_size": 2604868,
3   "maintenance_work_mem": 765255,
4   "max_wal_size": 2559064,
5   "max_worker_processes": 2,
6   "shared_buffers": 193629,
7   "temp_buffers": 6601,
8   "wal_buffers": 1696337,
9   "work_mem": 5947
10 }
```

Feature 2

...

```
1 {
2   "effective_cache_size": 3630833,
3   "maintenance_work_mem": 4034426,
4   "max_wal_size": 1998371,
5   "max_worker_processes": 6,
6   "shared_buffers": 2377061,
7   "temp_buffers": 57010,
8   "wal_buffers": 727571,
9   "work_mem": 37737
10 }
```

Feature n

Sampling

```
1 time(sec), throughput(req/sec),
2 0,940.864,4791.046,33.078,5002.
```

Average throughput (5min)

```
1 Transaction Type Index,Transaction Name,Start
2 4,Delivery,1678361810.799424,95955,5,0
3 1,NewOrder,1678361810.799584,57134,0,0
4 2,Payment,1678361810.799587,34646,9,0
5 1,NewOrder,1678361810.801177,56958,8,0
6 1,NewOrder,1678361810.801192,57178,7,0
7 2,Payment,1678361810.801196,33078,3,0
8 4,Delivery,1678361810.802719,96234,6,0
```

```
283198 2,Payment,1678362105.795620,4999754,9,0
283199 2,Payment,1678362105.796875,5002308,5,0
283200 2,Payment,1678362105.798137,5001561,4,0
283201 3,OrderStatus,1678362105.798137,4999451,2,0
```


Log of all transactions and their latency (5min)

Pre-processing

```
1 Transaction Type Index,Transaction Name,Start  
2 4,Delivery,1678361810.799424,95955,5,0  
3 1,NewOrder,1678361810.799584,57134,0,0  
4 2,Payment,1678361810.799587,34646,9,0  
5 1,NewOrder,1678361810.801177,56958,8,0  
6 1,NewOrder,1678361810.801192,57178,7,0  
7 2,Payment,1678361810.801196,33078,3,0  
8 4,Delivery,1678361810.802719,96234,6,0
```

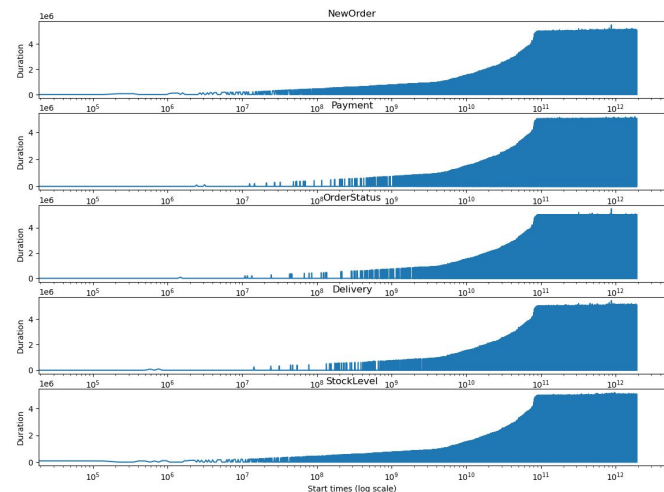

```
283198 2,Payment,1678362105.795620,4999754,9,0  
283199 2,Payment,1678362105.796875,5002308,5,0  
283200 2,Payment,1678362105.798137,5001561,4,0  
283201 3,OrderStatus,1678362105.798137,4999451,2,0
```

Log of all transactions and their latency (5min)



Pre-processing

$$JSD(A_{default} || B_{default})$$



5 JSD scores


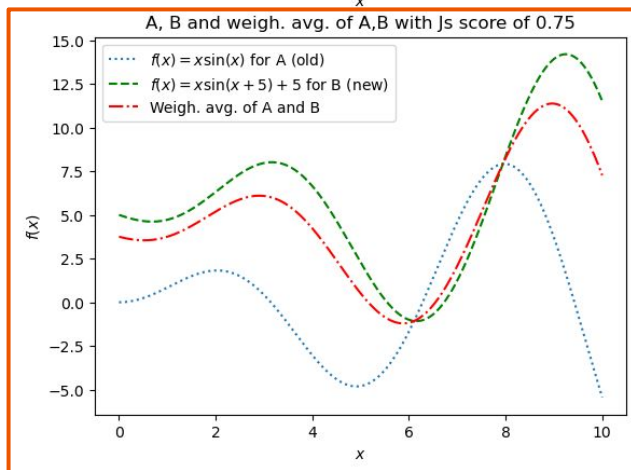
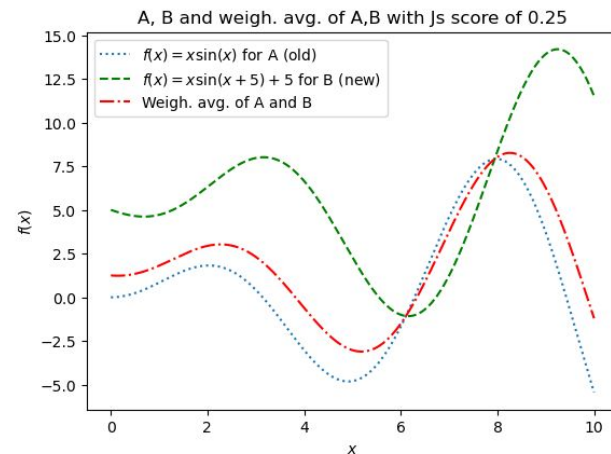
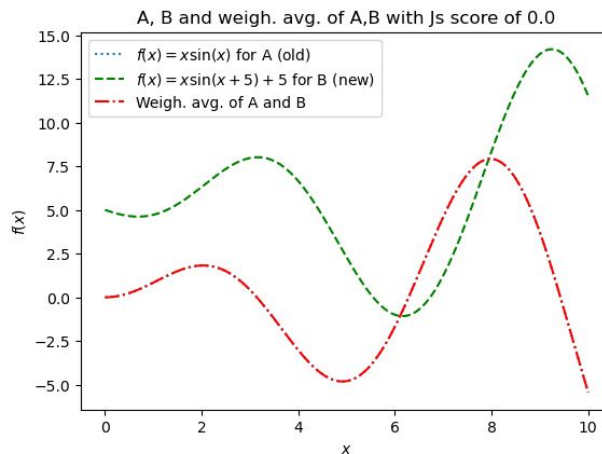


New ensemble model

$$n_{new} = f(n_{old}, \theta) = \theta * n_{old}$$

$$\hat{y} = \begin{cases} M_1(x_{old}) & ; \theta = 0 \\ M_2(x_{new}) & ; \theta = 1 \\ M_1(x_{old})(1 - \theta) + M_2(x_{new})\theta & ; 0 \leq \theta \leq 1 \end{cases}$$

New ensemble model



Results



Default value

A

1811 request/sec

B

1522 request/sec

$$\text{JSD} = 0.33$$



Sampling

A

100 samples

9 hours

B

33 samples

3 hours

$$\text{JSD} = 0.33$$



Ensemble prediction for B

Default: 1522 req/sec

Best config: 1915 req/sec

JSD = 0.33

25% faster

In only $\frac{1}{3}$ of the tuning time



Evaluation

Contributions

Better than guessing

Faster than default

Reduced tuning bottleneck!

Limitations

Inconsistent results

Not enough data for hypothesis test yet

Proof of concept but data too inconsistent



Project management

Well conceived project

- All necessary research, analysis and design work completed
- Agile development approach

Could have received code base earlier from (Barbulescu, Triantafillou; 2022)

- Progress no always consistent with specification report

Unforeseen problems well detected and overcome

- DB Caching issues
- Legacy version of DB - CLI problems

Q & A